


CLOUD HASKELL

Haskell support for the development of distributed
applications

Alexandra Flinta
George Abordioaie

Distributed applications

- ▶ Distributed applications (distributed apps) are applications or software that runs on multiple computers within a network at the same time and can be stored on servers or with cloud computing.



Distributed applications

- ▶ Can communicate with multiple servers or devices on the same network from any geographical location
- ▶ Are broken up into two separate programs: the client software and the server software:
 - ▶ The client software or computer accesses the data from the server or cloud environment
 - ▶ The server or cloud processes the data
- ▶ It can failover to another component if a distributed application component goes down



Distributed applications - benefits

- ▶ Allow multiple users to access the apps at once
- ▶ Developers, IT professionals or enterprises choose to store distributed apps in the cloud:
 - ▶ cloud elasticity
 - ▶ cloud scalability
 - ▶ ability to handle large applications
 - ▶ ability to handle workloads



Cloud Haskell – Erlang-style concurrency

- ▶ Cloud Haskell is a set of libraries that bring Erlang-style concurrency and distribution to Haskell programs
 - ▶ Fast process creation/destruction
 - ▶ Ability to support >> 10 000 concurrent processes with largely unchanged characteristics.
 - ▶ Fast asynchronous message passing.
 - ▶ Copying message-passing semantics (share-nothing concurrency).
 - ▶ Process monitoring.
 - ▶ Selective message reception.



Cloud Haskell – New approach

- ▶ Has be re-written from the ground up and supports a rich and growing number of features for:
 - ▶ building concurrent applications using asynchronous message passing
 - ▶ building distributed computing applications
 - ▶ building fault tolerant systems
 - ▶ running Cloud Haskell nodes on various network transports
 - ▶ working with several network transport implementations (and more in the pipeline)
 - ▶ supporting *static* values (required for remote communication)



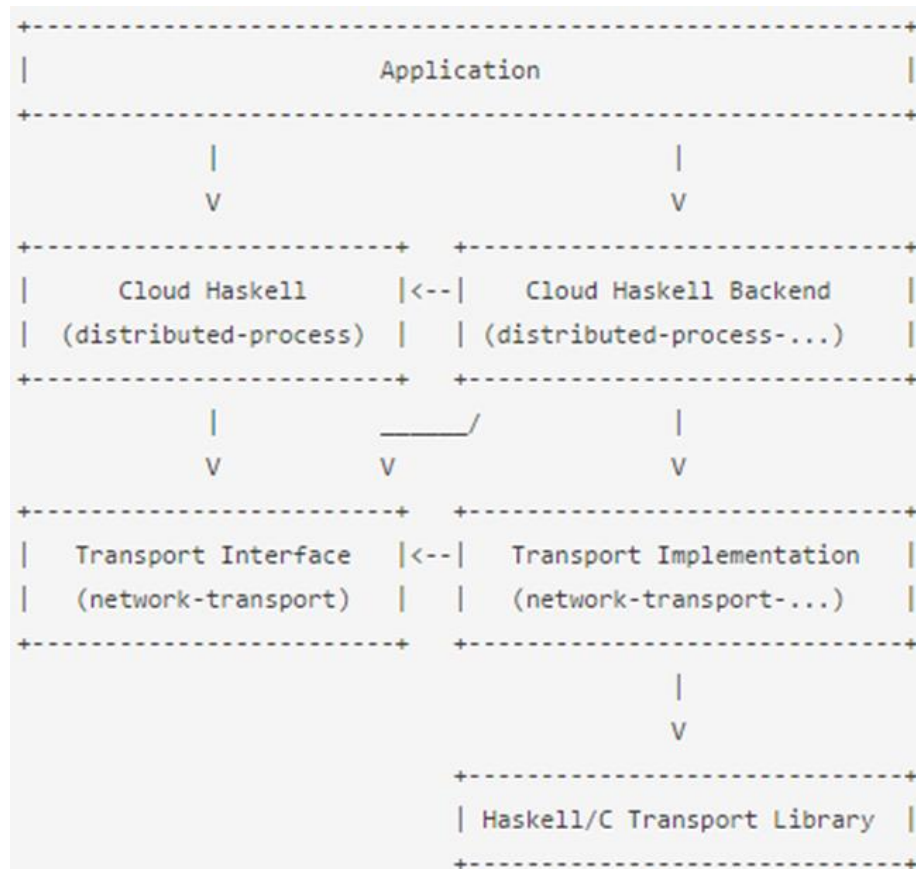
Cloud Haskell – Components

- ▶ Cloud Haskell comprises the following components, some of which are complete, others experimental.
 - ▶ [distributed-process](#): Base concurrency and distribution support
 - ▶ [distributed-process-platform](#): The Cloud Haskell Platform - APIs
 - ▶ [distributed-static](#): Support for static values
 - ▶ [rank1dynamic](#): Like Data.Dynamic and Data.Typeable but supporting polymorphic values
 - ▶ [network-transport](#): Generic Network.Transport API
 - ▶ [network-transport-tcp](#): TCP realisation of Network.Transport
 - ▶ [network-transport-inmemory](#): In-memory realisation of Network.Transport (incomplete)
 - ▶ [network-transport-composed](#): Compose two transports (very preliminary)
 - ▶ [distributed-process-simplelocalnet](#): Simple backend for local networks
 - ▶ [distributed-process-azure](#): Azure backend for Cloud Haskell (proof of concept)



Cloud Haskell – Support for Haskell

The following diagram shows dependencies between the various subsystems, in an application using Cloud Haskell, where arrows represent explicit directional dependencies.



Cloud Haskell – Support for Haskell

In this diagram, the various nodes roughly correspond to specific modules:

Cloud Haskell	: <code>Control.Distributed.Process</code>
Cloud Haskell	: <code>Control.Distributed.Process.*</code>
Transport Interface	: <code>Network.Transport</code>
Transport Implementation	: <code>Network.Transport.*</code>

- ▶ *Control.Distributed.Process* module defines abstractions such as nodes and processes
- ▶ The Transport interface provided by the *Network.Transport* module is used by the Cloud Haskell interface and backend
- ▶ The *Network.Transport.** module provides a specific implementation for the current transport



Concurrency and distributed applications

- ▶ The *Process Layer* is where Cloud Haskell's support for concurrency and distributed programming are exposed to application developers.
- ▶ Processes reside on nodes, which in our implementation map directly to the `Control.Distributed.Processes.Node` module. Given a configured `Network.Transport` backend, starting a new node is fairly simple:

```
newLocalNode :: Transport -> IO LocalNode
```

- ▶ Given a new node, there are two primitives for starting a new process.

```
forkProcess :: LocalNode -> Process () -> IO ProcessId  
runProcess  :: LocalNode -> Process () -> IO ()
```



Building Examples

▶ Prerequisites:

- ▶ Haskell environment
- ▶ Stack: cross-platform program for developing Haskell projects
- ▶ Cabal: system for building and packaging Haskell libraries
- ▶ `distributed-process` and `network-transport-tcp` libraries from Cloud Haskell Platform



Building the libraries

- ▶ Download and install Cabal
- ▶ Download and unzip `distributed-process` and `network-transport-tcp` libraries
- ▶ Go to the root folder of every library and type: `cabal install library-name.cabal`



Setting up the project

- ▶ Setting up the project is made using stack
- ▶ Type `stack new project-name` in a fresh new directory. This will populate the directory with a number of files, chiefly *stack.yaml* and **.cabal* metadata files for the project.
- ▶ You'll want to add distributed-process and network-transport-tcp to the build-depends stanza of the executable section



Create a node

- ▶ Cloud Haskell's lightweight processes reside on a “node”, which must be initialized with a network transport implementation and a remote table.

```
import Network.Transport.TCP (createTransport, defaultTCPParameters)
import Control.Distributed.Process
import Control.Distributed.Process.Node
```

```
main :: IO ()
main = do
    Right t <- createTransport "127.0.0.1" "10501" defaultTCPParameters
    node <- newLocalNode t initRemoteTable
    ....
```



Sending Messages

- ▶ We start a process by evaluating `runProcess` which takes a `Process` and a node to run:

```
-- in main
_ <- runProcess node $ do
  -- get our own process id
  self <- getSelfPid
  send self "hello"
  hello <- expect :: Process String
  liftIO $ putStrLn hello
  return ()
```

- ▶ Each process has an identifier associated to it.
- ▶ Each process also has a mailbox associated with it. Messages sent to a process are queued in this mailbox in an asynchronous .



Sending Messages II

- ▶ `receiveWait` and similarly named functions can be used with the `Match` data type to provide a range of advanced message processing capabilities.

```
-- Test our matches in order against each message in the queue  
receiveWait [match logMessage, match replyBack]
```



Sending Messages III

- ▶ In the `echo server`, our first match prints out whatever string it receives. If the first message in our mailbox is not a `String`, then our second match is evaluated.

```
-- Spawn another worker on the local node  
echoPid <- spawnLocal $ forever $ do
```

```
replyBack :: (ProcessId, String) -> Process ()  
replyBack (sender, msg) = send sender msg
```



Serializable Data

- ▶ Processes may send any type of data as long as they implement the `Serializable` typeclass:

```
class (Binary a, Typeable) => Serializable a
instance (Binary a, Typeable a) => Serializable a
```

- ▶ Any type that is `Binary` and `Typeable` is `Serializable`.
- ▶ For custom data types a `Typeable` instance is given by the compiler and a `Binary` instance can be generated as such:

```
{-# LANGUAGE DeriveDataTypeable #-}
{-# LANGUAGE DeriveGeneric #-}

data T = T Int Char deriving (Generic, Typeable)

instance Binary T
```



Spawning Remote Processes

- ▶ The behaviour of processes is determined by an action in the `Process` monad.
- ▶ However, actions in the `Process` monad, no more serializable than actions in the `IO` monad. If we can't serialize actions, then how can we spawn processes on remote nodes?
- ▶ Solution? – Static actions ...



Spawning Remote Processes II

- ▶ Static actions are not easy to construct by hand, but fortunately Cloud Haskell provides a little bit of Template Haskell to help. If: `f :: T1 -> T2`

- ▶ Then: `$(mkClosure 'f) :: T1 -> Closure T2`

- ▶ You can turn any top-level unary function into a `Closure` using `mkClosure`

- ▶ For curried functions you need to move the arguments into tuples and make a mapping into the remote table.

```
sampleTask :: (TimeInterval, String) -> Process ()
sampleTask (t, s) = sleep t >> say s

remotable ['sampleTask]
```



Spawning Remote Processes III

- ▶ The call to `remotable` implicitly generates a remote table by inserting a top-level definition `__remoteTable :: RemoteTable -> RemoteTable` in our module for us. We compose this with other remote tables in order to come up with a final, merged remote table for all modules in our program.

